

COMP 341 Intro to AI

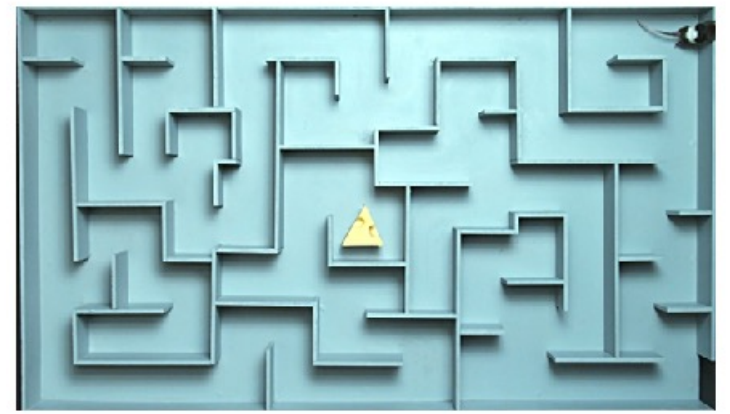
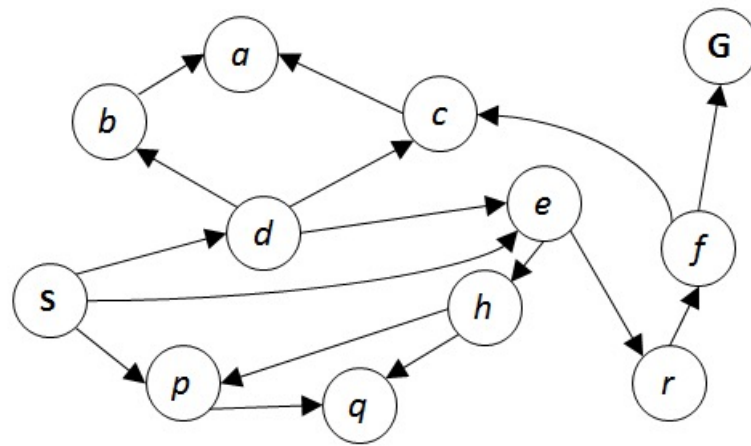
Local Search



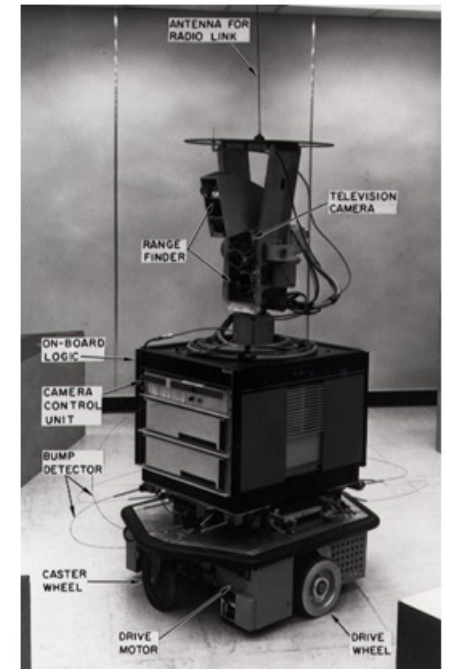
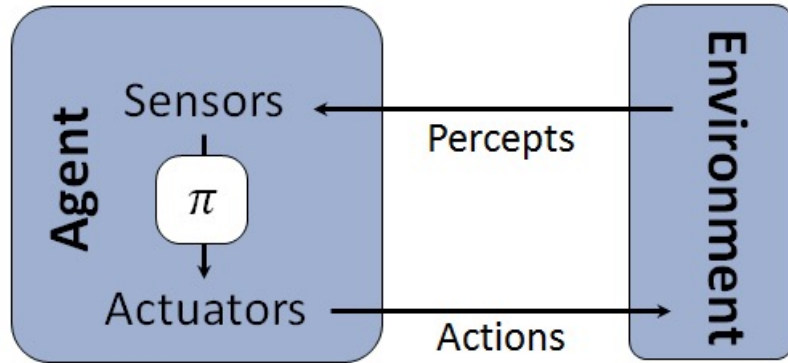
Asst. Prof. Barış Akgün
Koç University

Today

- Recap
- Local Search
 - “Hill-Climbing”
 - Simulated Annealing
 - Local Beam Search
 - Genetic Algorithms
 - Gradient Ascent/Descent
- Online Search
- Beyond Classical Search: Some Topics



Previously on Intro to AI

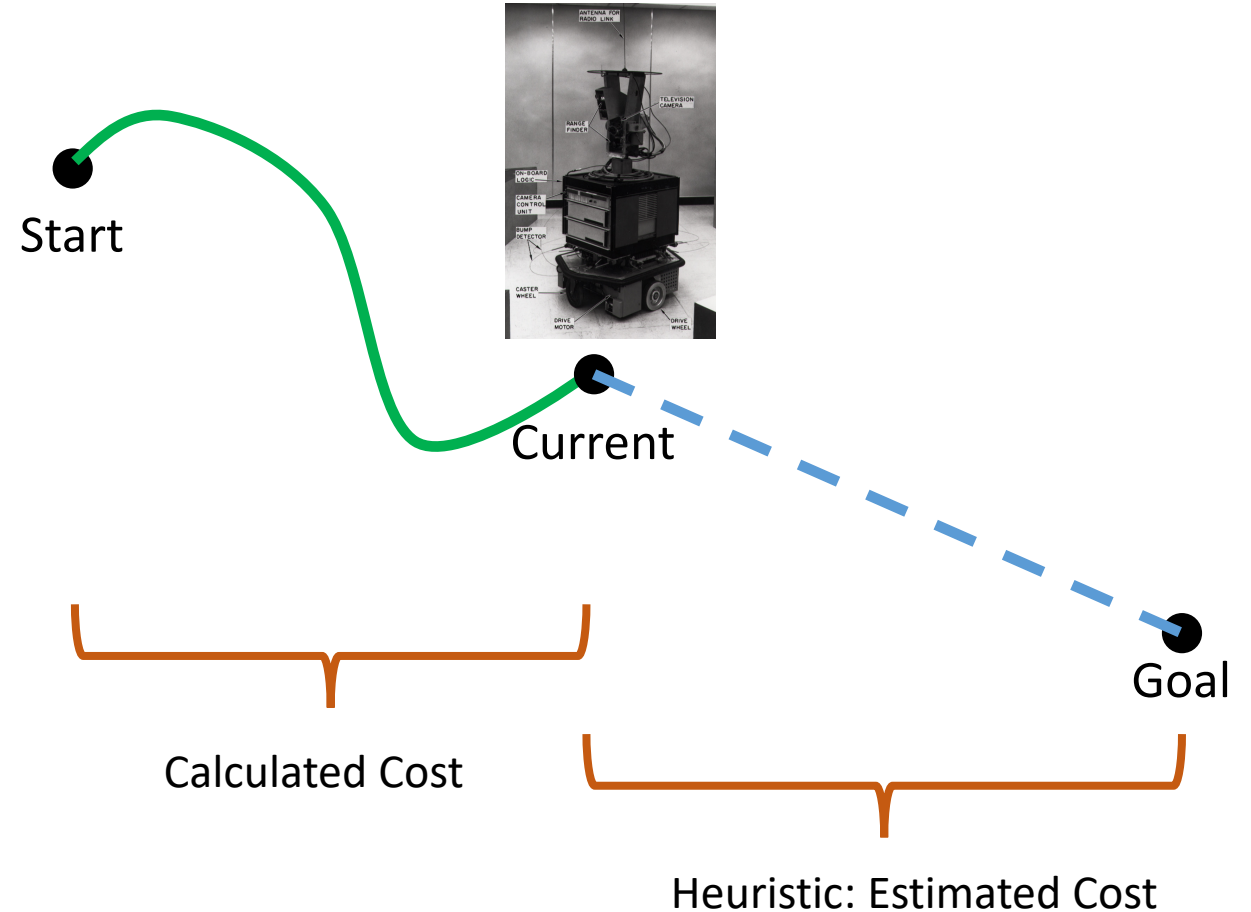


Search

- Uninformed
 - DFS, BFS, UCS
 - No domain knowledge
- Informed
 - Greedy, A*
 - Heuristics based on domain/problem knowledge
- Solution is a path to goal

Heuristics

- How promising is a given state?
 - Admissibility/Consistency
 - Dominance
- Design:
 - Solution Costs to Relaxed Problems (e.g., 8 puzzle)
 - Geometric Limits (e.g., Euclidean/Manhattan)
 - Creativity 😊



Local Search



Local Search

- Classical Search:
 - Solution is path to a goal state
- Local Search:
 - Solution is the goal state itself
 - Search for a solution when the path doesn't matter
- Optimization: Get to a “better” state (best state if possible!)


Local Search Applications

- Integrated circuit design
- Factory floor layout
- Scheduling
- Routing
- Portfolio Management
- Network optimization
- ...

Local Search

- Formulation:
 - Current State
 - Transition/Successor Function
 - Evaluation Function
- Algorithms: Move towards Better States
 - Complete: Find a solution if one exists
 - Optimal: Find the best state
- New Concept: State Space “Landscape”
- Usually easy to code!


1D State Space Example



Hill-Climbing (in Heavy Fog with Amnesia)


```
function HILL-CLIMBING(problem) returns a state that is a local maximum  
  current  $\leftarrow$  MAKE-NODE(problem.INITIAL-STATE)  
  loop do  
    neighbor  $\leftarrow$  a highest-valued successor of current  
    if neighbor.VALUE  $\leq$  current.VALUE then return current.STATE  
    current  $\leftarrow$  neighbor
```

N-Queens





Place N-Queens on a NxN chessboard so that no queen attacks another. (A queen can attack any piece on the same column, row or diagonal.)

- State?
 - Position per column
- Successor Function?
 - Move a piece in a column
- Evaluation Function?
 - Number of attacks




$h = 5$



$h = 2$

Example: 8-Queens




18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	13	16	13	16	16
14	17	15	14	16	16	16	16
17	16	18	15	15	14	15	16
18	14	15	15	14	16	16	16
14	14	13	17	12	14	12	18

Current evaluation: 17
(all attacks on the same line are counted)

Hill Climbing Properties

- Complete
 - No
- Optimal
 - No
- Time Complexity
 - $O(d)$ – d : longest path to solution (could be infinite!)
- Space Complexity
 - Constant



Drawbacks:

- Plateaus
- Local Minima
- No memory

Hill Climbing – Avoiding Drawbacks?

- Get out of Plateaus:
 - Random walk among equally valued states – can escape “shoulders”
 - Infinite Loop, e.g. local flat maxima: Move thresholding
- Avoid Local Minima:
 - Random restart: Start the problem at different states after termination
 - Probabilistically complete but not efficient
- Add Memory:
 - Hill climbing from multiple states in parallel
 - Exchange information

Some Randomized Variants


- Stochastic HC: choose randomly from uphill moves
- First choice Stochastic HC: generate successors randomly, pick 1st uphill
- Random Restart HC: try and try HC until a goal is found
- What is left?
 - Maybe sometimes move downhill?

Reality Check for HC

- HC is too greedy
 - Never moving downhill is incomplete
 - Real problems have lots of local maxima
- Randomness can bring completeness, but it is inefficient
- So?
 - Combine them!

Simulated Annealing

- Idea:
 - Sometimes move downhill to escape local maxima
 - i.e. Combine random walk with hill climbing
 - Gradually decrease the size and frequency of the random walk
- T:
 - Probability of picking a random neighbor instead of best
 - “Temperature parameter”
- Slowly decrease the temperature – hence annealing!



Simulated Annealing

function SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state

#schedule, a mapping from time to “temperature”

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

while true **do**

$T \leftarrow \text{schedule}(t)$

if $T = 0$ **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow \text{next}.\text{VALUE} - \text{current}.\text{VALUE}$

if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{\Delta E/T}$

reduce T over time

Select a random neighbor and
calculate the change in value

If the change is positive, move to
the next state
else move with some probability

If T is decreased “slow enough”, then SA is probabilistically complete and globally optimal!

SA Temperature Scheduling Examples

$$T(t) = \frac{d}{\log(t)}, d > 0, t > 1$$


$$T(t) = \alpha T(t - 1), 1 > \alpha > 0$$

$$T(t) = T(t - 1) - k, k > 0$$


Local Beam Search - Memory

- **Idea:** Keep k states instead of 1 and chose top k of all their successors
 - Not the same as k searches run in parallel! Searches that lead to good states recruit other searches to join them


Local Beam Search - Memory



Local Beam Search - Memory



Local Beam Search - Memory



Local Beam Search - Memory

- **Idea:** Keep k states instead of 1 and chose top k of all of their successors
 - Not the same as k searches run in parallel! Searches that and good states recruit other searches to join them
- **Problem:** quite often, all k states end up on same local hill
- **Idea 2:** choose k successors randomly, biased towards good ones
 - Stochastic Beam Search
 - Natural selection anyone?

Genetic Algorithms

- Stochastic local beam search + generate successors from **pairs** of states
- Formulation
 - state: *finite alphabet*
 - k states: *population*
 - Evaluation function called *fitness function*
 - Successors from two parents

Brainstorming

- Let's formulate the below problems:
 - State
 - Fitness Function
 - Successor Function (?)
- Evolve a binary string of length n that only contains 1
- Evolve a symmetric binary string

Genetic Algorithms

24748552
32752411
24415124
32543213

(a)

Initial Population

states encoded from a finite alphabet
k initial states, usually randomly selected


Genetic Algorithms

24748552	24
32752411	23
24415124	20
32543213	11

Each state has a “fitness” which is given by the fitness/evaluation function

(a)	(b)
Initial Population	Fitness Function


Genetic Algorithms



Select pairs of states for **reproduction**, weighted by their fitness

This state was not lucky enough

Genetic Algorithms



Reproduction: randomly choose a cross-over point, create two new states by mixing the pair

Genetic Algorithms




Mutation: these new states are also subject to random mutations of single digits in the state

Genetic Algorithms – 8 Queens

- 8 Queens Puzzle: placing eight chess queens on an 8×8 chessboard so that no two queens threaten each other
- Formulate:
 - State?
 - Fitness Function?
 - Do we need a transition function?

Genetic Algorithms – 8 Queens

- 8 Queens Puzzle: placing eight chess queens on an 8×8 chessboard so that no two queens threaten each other



Local Search Summary

- Hill Climbing
 - Simulated Annealing
 - Genetic Algorithms
 - There are more (Particle Swarm Optimization is/was very popular)
-
- These can be applied to continuous state spaces (how?)
 - These are all “derivative free algorithms”

Derivative Based: Gradient Descent/Ascent

- State x : multivariate and continuous
- Objective function $f(x)$: Differentiable around x
- **Idea**: Move x in the direction of decreasing/increasing f
- **How**: Derivatives!
- Need to calculate the **gradient**:

Getting into the realm of optimization!

$$\nabla f(x) = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right)$$

Can be computed *analytically* or *numerically*


Gradient Descent/Ascent

- Implementation:

While SomeCondition:

$$x(t + 1) = x(t) \pm \alpha \nabla f(x)$$

- SomeCondition:
 - Maximum iterations
 - $|x(t+1) - x(t)| < \text{small}$
 - ...
- More complicated versions exist such as calculating higher order derivatives (e.g. *Newton-Raphson*), adaptive step size (e.g. momentum) etc.



A few Words on Gradient Descent/Ascent

- This family of algorithms is used in **a lot of** applications!
- Very simple to code the base version
- Works in any number of dimensions
 - Even in infinite-dimensions! Just need to change the derivative
- Inefficient
- Some state-space landscapes wreak havoc
- In any case, a good first tool to tackle an applicable problem!

Online Search

- **Offline:** simulate the world and reason about a plan to get to a goal
- **Online:** solve the search problem while executing actions
- Interleaves search and execution



Remaining parts of this slide deck will not be in the exams. Similar topics introduced later will be.

Why Online Search?

- **Dynamic environment**, things change too quickly to plan
- **Nondeterministic environment**, deal with what happens rather than planning for all contingencies
- **Hard to model**, cannot get a reasonable model to do search on
- **Problem**: a shortsighted view may lead the agent astray


Online Search

- $Actions(s)$ = actions allowed in state s
- $c(s, a, s')$ = cost of action going from state s to state s' with action a
- $Goal-test(s)$
- No transition model! Do things and see what happens
 - You know s' when you get there!
- Issues:
 - Irreversible actions
 - Dead-ends
 - Is environment **safe to explore?**

Online Searching with Algos Learned so Far

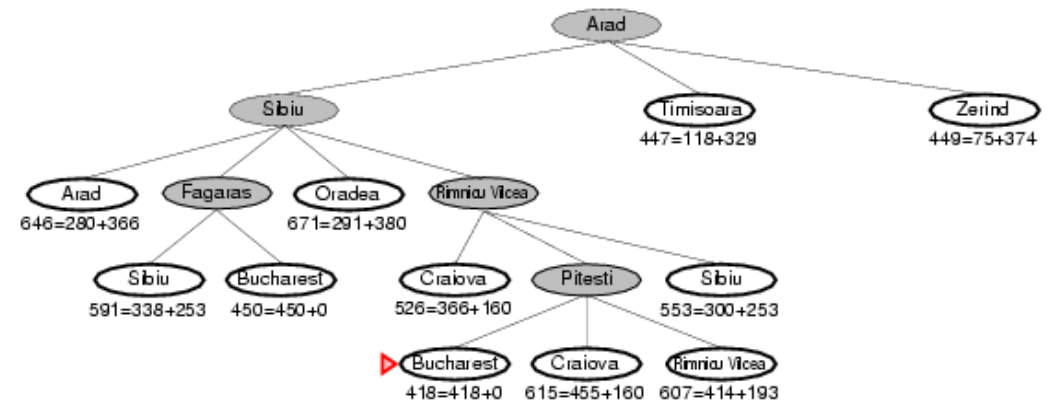
- Hill-Climbing?

- Yes!
- In fact, most local search methods
- Obviously, no random restarts



- A*?

- No
- Cannot jump around the state space!



But I really like A*!

- Learning Real-Time A*
- Follow $f(n)$ locally but note that initial state does not matter!
- Learning: update $h(s)$ with experience to keep from getting stuck in local minima
- If interested read:

Richard E. Korf, *Real-time heuristic search*, Artificial Intelligence 1990

Beyond Classical Search

- What if actions are non-deterministic?
 - Contingency Plans, plan for multiple outcomes
- What if percepts do not give a whole state? (Partially observable)
 - Reason about **belief states**
 - Things get more complex
 - The full version is not in the scope of this course