

COMP 350

Selected Topics - Introduction to DevOps



Lecture 1

Introduction

Hakan Ayrar

Administrivia

- Textbooks
 - first half of the course is about Linux system administration:
 - The Linux Command Line, William E. Shotts, No Starch Press
 - on the second half we will focus on actual DevOps
 - mainly online references (git, docker,...) and lecture notes
 - but also Practical DevOps, Joakim Verona, Packt Publishing
- Office Hours: Wednesday 14:20-15:20 or any other time by appointment
- Participation and attendance is 10% of course grade
- You are expected to do some reading and practicing
- Grading
 - 2x Projects 30%
 - Midterm Exam 25%
 - Final Exam 35%
 - Participation and attendance 10%



Syllabus

Meeting Times	Subject
1	Linux Operating System, Kernel and kernel modules, Kernel boot parameters, Boot process of Linux. Initial ram disk, systemd/initStorage Management and File Systems on Linux. • Common file systems • Access Control for files and directories, permissions
2	Devices, Block devices, mounting block devices for storage • Create and manage basic volumes, spanning volumes, redundancy, volume management and RAID configurations • Swap partition, virtual memory, Swap vs ZSwap vs zRAM
3	What's a process? How to manage processes in Linux? • Monitoring processes • StdIO, StdErr, I/O Redirection, Pipes, Filtering with grep
4	Shell as a process, CLI interaction and scripting • Terminal multiplexers • Secure shell remote access • User management, access control and System Users • cron • System Logs
5	Package management systems and scripting/programming ecosystems: • Apt, yum, deb, rpm • Python, pip, conda • Nodejs and npm • Ruby and gemsBuild tools: make, cmake, GNU Build System (aka autotools), Apache Ant/Maven, Gradle
6	Code repositories, Versioning, Git and github in depth: repo, branch, commit, clone, fast forward, rebase, cherry picking
7	Container systems, virtualization and scaling: • Quick look at the beginnings: Chroot jail • Virtualization • Containerization • Docker, docker file, docker compose
8	Networking in virtualized/containerized environments • NAT, Port forwarding, Routing, Packet filtering
9	Container Orchestration: Docker swarm, Kubernetes, Amazon Elastic Container Service (ECS) • Load balancing for scaling

Syllabus (cont'd)

10	Microservice Architecture: philosophy, building APIs, REST APIs, SOAP, gRPC, XML, YAML
11	Cloud service providers: Amazon Web Services, Microsoft Azure, Google Cloud Platform Serverless: AWS Lambda, Google Cloud Functions, Azure Serverless
12	Infrastructure as Code (IaC): Terraform, OpenTF, AWS CloudFormation,
13	Continuous Integration (CI)/Continuous Delivery (CD) : • Gitlab CI, Azure pipelines, AWS Code pipeline • Jenkins, Travis, Teamcity, Circle CI
14	Configuration Management: Chef, Puppet, Salt, Ansible, Terraform, CFEngine

What are we interested in?

- Linux Operating System
- System Administration
- Storage Management and file systems
- Inter-process communications

System Admin

- Computer communications and Networking

Network Admin

- Well-structured documents
- Regular expressions
- Computer Languages
- Scripting/Programming platforms and related package management systems

Developer

- Container systems, virtualization and scaling for deployment, and production environment
- Dev Ops best practices
- Code repositories, Versioning, automated build and Continuous Integration (CI)
- Compile arbitrary software directly from source

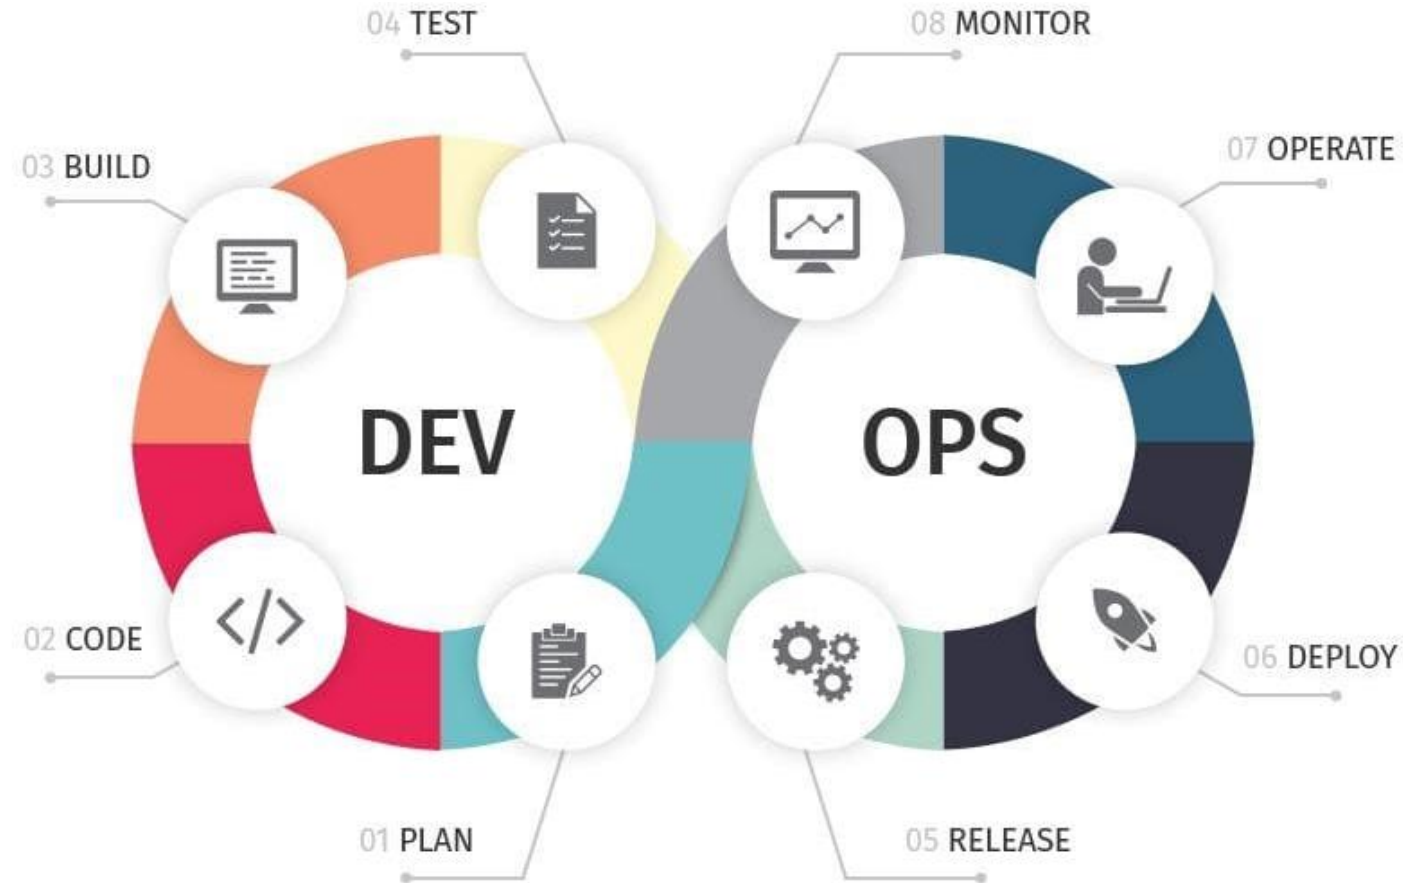
Dev Ops

What is DevOps?

- DevOps is a set of practices that combines software development (Dev) and IT operations (Ops).
- It aims to shorten the systems development life cycle and provide continuous delivery with high software quality.
- DevOps is complementary with Agile software development.
- As DevOps is intended to be a cross-functional mode of working, those who practice the methodology use different sets of tools—referred to as "toolchains"—rather than a single one.
- These toolchains are expected to fit into one or more of the following categories, reflective of key aspects of the development and delivery process:
 1. **Coding** – code development and review, source code management tools, code merging.
 2. **Building** – continuous integration tools, build status.
 3. **Testing** – continuous testing tools that provide quick and timely feedback on business risks.
 4. **Packaging** – artifact repository, application pre-deployment staging.
 5. **Releasing** – change management, release approvals, release automation.
 6. **Configuring** – infrastructure configuration and management, infrastructure as code tools.
 7. **Monitoring** – applications performance monitoring, end-user experience.
- Some categories are more essential than others.

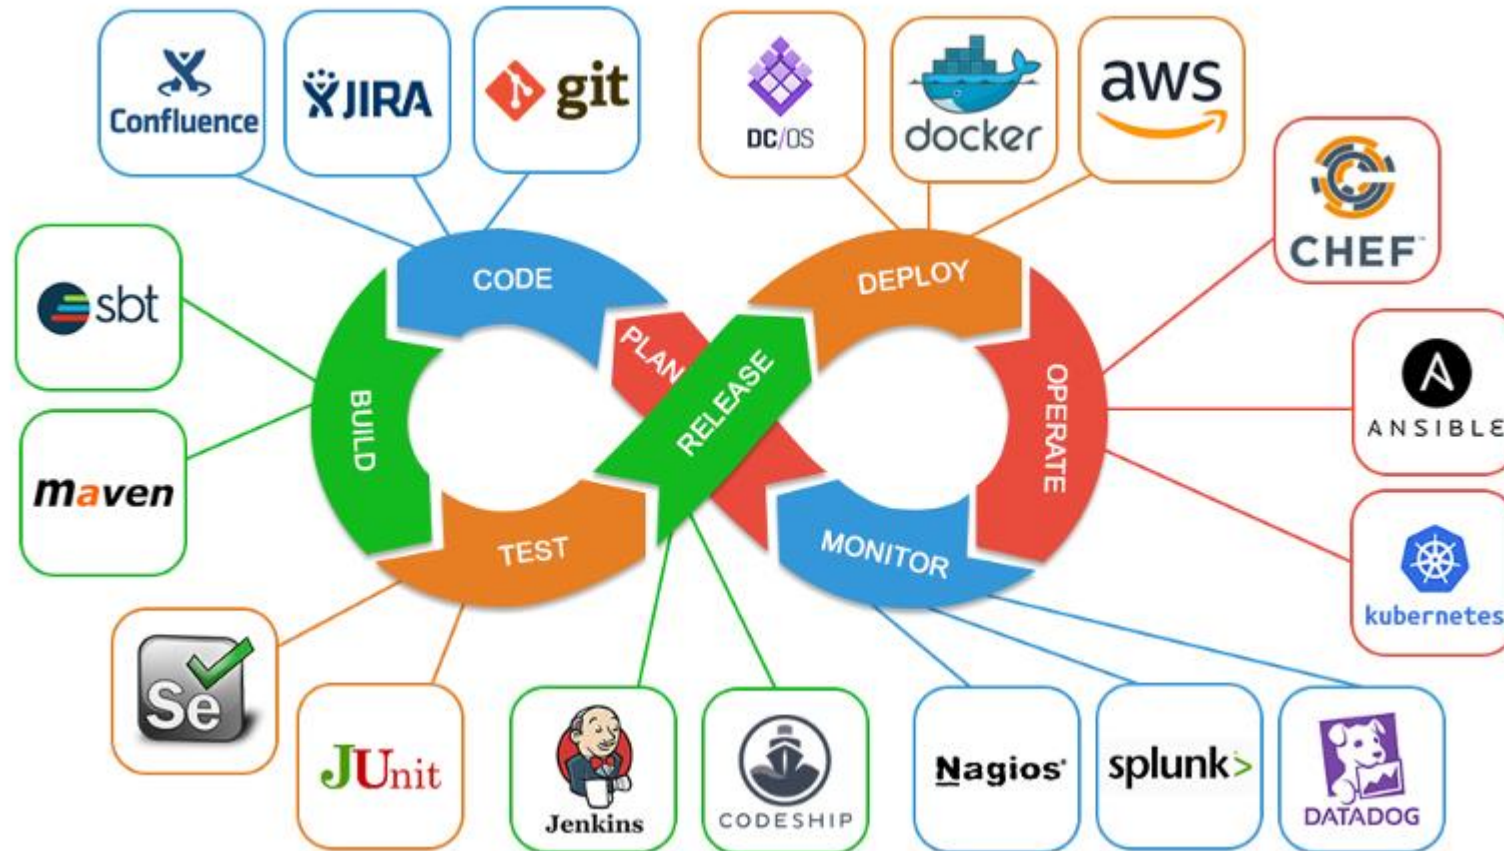
DevOps Cycle

- DevOps automates the **building** of both infrastructure and applications.
- DevOps applies a series of automated tests to both infrastructure and applications (unit, integration, performance etc) to prove they meet their functional and non-functional requirements.
- When proven to be fit to release, DevOps automates the **deployment** to end-users.
- DevOps primarily applies to businesses that build or assemble their own software, versus just using pre-built software.
- Forsgren et al. found that IT performance is strongly correlated with DevOps practices like source code management and continuous **delivery**. (Nicole Forsgren; Gene Kim; Nigel Kersten; Jez Humble (2014). ["2014 State of DevOps Report"](#))



Some tools used on different phases of the cycle

- If you are highly proficient on any 3 of the tools below you can start working on a big company tomorrow.



What is the goal of DevOps?

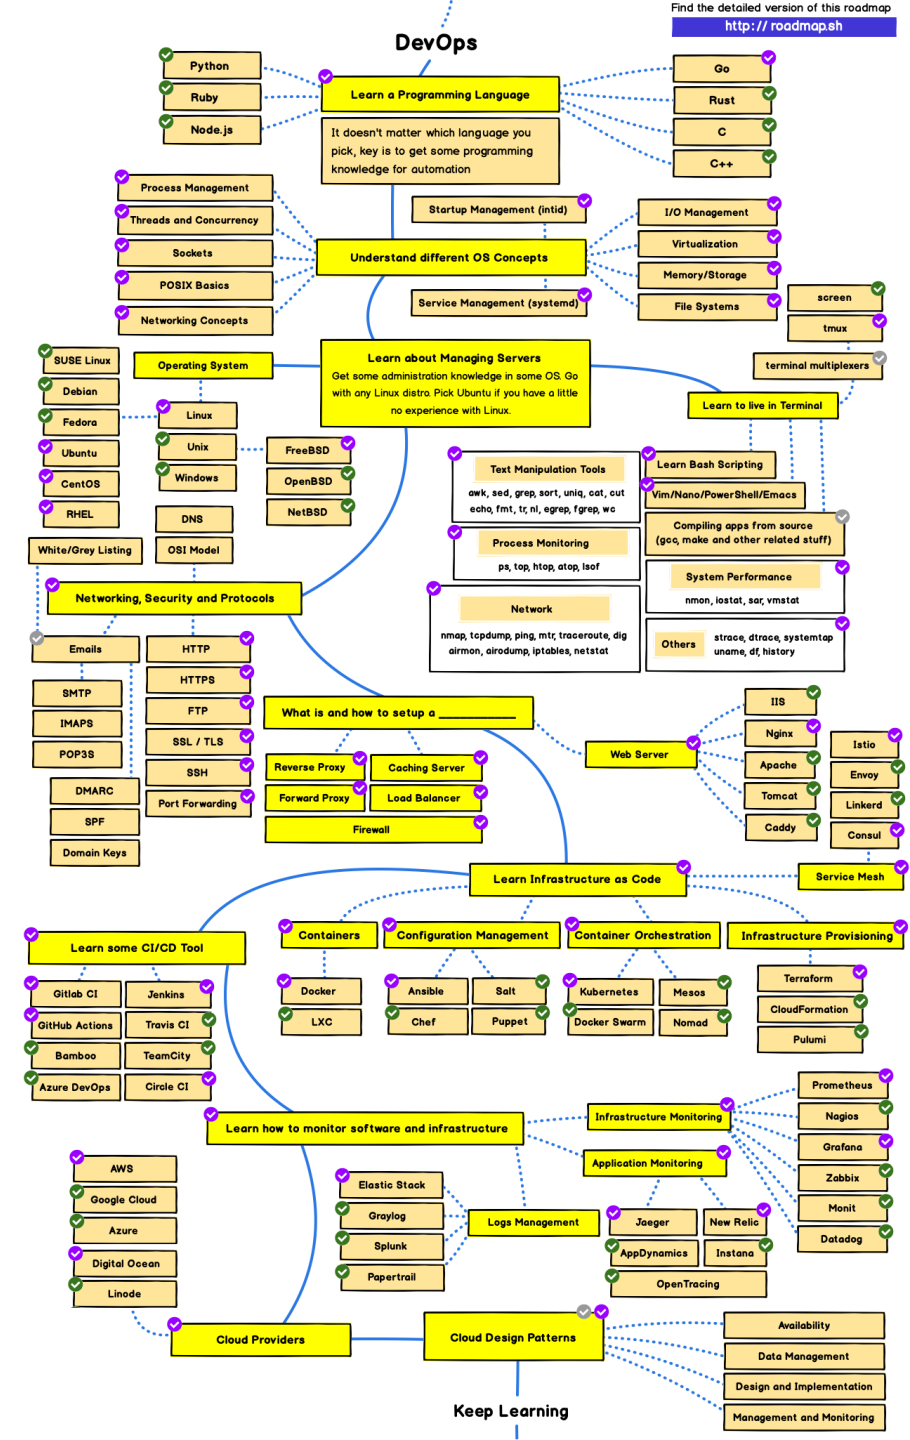
- The goals of DevOps span the entire delivery pipeline. They include:
 - Improved deployment frequency
 - Faster time to market
 - Lower failure rate of new releases
 - Shortened lead time between fixes
 - Faster mean time to recovery (in the event of a new release crashing or otherwise disabling the current system)
- DevOps as a job title: While DevOps describes an approach to work rather than a distinct role (like system administrator), job advertisements are increasingly using terms like "DevOps Engineer".

Aim of the course

- We will spend the first half of the term for familiarization with fundamental system administration tasks for enterprise environment.
- The purpose of this course is **not** to make a full-time **system admin** out of you.
- The purpose is to introduce you the DevOps tools and techniques needed for a **developer** to manage and orchestrate the IT operations needed during the steps of the DevOps cycle.
 - DevOps Engineer = Developer + some Sys Admin skills + extra skills specific to DevOps
- Obviously the length of a semester is not long enough to cover every tool and technique in detail.
- You are expected to do a lot of reading and hands-on practice.

An Unofficial Roadmap

1. Learn a Programming Language
2. Understand different OS concepts
3. Learn to Live in terminal
4. Networking and Security
5. What is ... and how to setup? (i.e. app servers)
6. Learn Infrastructure as code
7. Learn some Continuous Integration and Delivery (CI/CD) tools
8. Learn to monitor software and infrastructure
9. Learn about Cloud Providers



The Linux Command Line

PART 1 - LEARNING THE SHELL

- WHAT IS THE SHELL?
- NAVIGATION
- EXPLORING THE SYSTEM
- MANIPULATING FILES AND DIRECTORIES
- WORKING WITH COMMANDS
- REDIRECTION
- SEEING THE WORLD AS THE SHELL SEES IT
- ADVANCED KEYBOARD TRICKS
- PERMISSIONS
- PROCESSES

PART 2 - CONFIGURATION AND THE ENVIRONMENT


- THE ENVIRONMENT
- A GENTLE INTRODUCTION TO VI
- CUSTOMIZING THE PROMPT

PART 3 - COMMON TASKS AND ESSENTIAL TOOLS

- PACKAGE MANAGEMENT
- STORAGE MEDIA
- NETWORKING
- SEARCHING FOR FILES
- ARCHIVING AND BACKUP
- REGULAR EXPRESSIONS
- TEXT PROCESSING
- FORMATTING OUTPUT
- PRINTING
- COMPILING PROGRAMS

PART 4 - WRITING SHELL SCRIPTS

...



What is Linux?

- A family of open-source Unix-like operating systems.
- It is based on the Linux kernel (an operating system kernel first released on September 17, 1991, by Linus Torvalds)
- Linux is typically packaged in a Linux distribution.
- Distributions include the Linux kernel and supporting system software and libraries.
- Desktop Linux distributions include a windowing system such as X11 or Wayland, and a desktop environment such as GNOME or KDE Plasma.
- Distributions intended for servers may omit graphics altogether, or include a solution stack such as LAMP (=Linux+Apache+MySQL+Php).
- Linux is freely redistributable, therefore anyone may create a distribution for any purpose.
- Without the mastery on OS you can neither be a good developer, nor a good sys admin (not to mention DevOps)

How widespread is Linux? (aka why is this course based on linux?)

- Because of the dominance of the Linux-based Android on smartphones, Linux has the largest installed base of all general-purpose operating systems.
- But it is used by only around **2.3%** of desktop computers.
- Linux is the leading operating system on servers (over **96.4%** of the top 1 million web servers' OS is Linux)
- It is also the only OS used on TOP500 supercomputers (since November 2017)
- Linux also runs on embedded systems (→ devices whose OS is typically built into the firmware)
 - This includes routers, automation controls, smart home technology (e.g. Google Nest), televisions, automobiles (e.g. Tesla, Audi, Mercedes-Benz, Hyundai, and Toyota), digital video recorders, video game consoles, and smartwatches.
 - The Falcon 9 space launch vehicle's and the Dragon 2 space craft's avionics use a customized version of Linux.
- **90%** of all cloud infrastructure is powered by Linux including supercomputers and cloud providers.
- **74%** of all smartphones in the world are Linux-based. (due to Android)

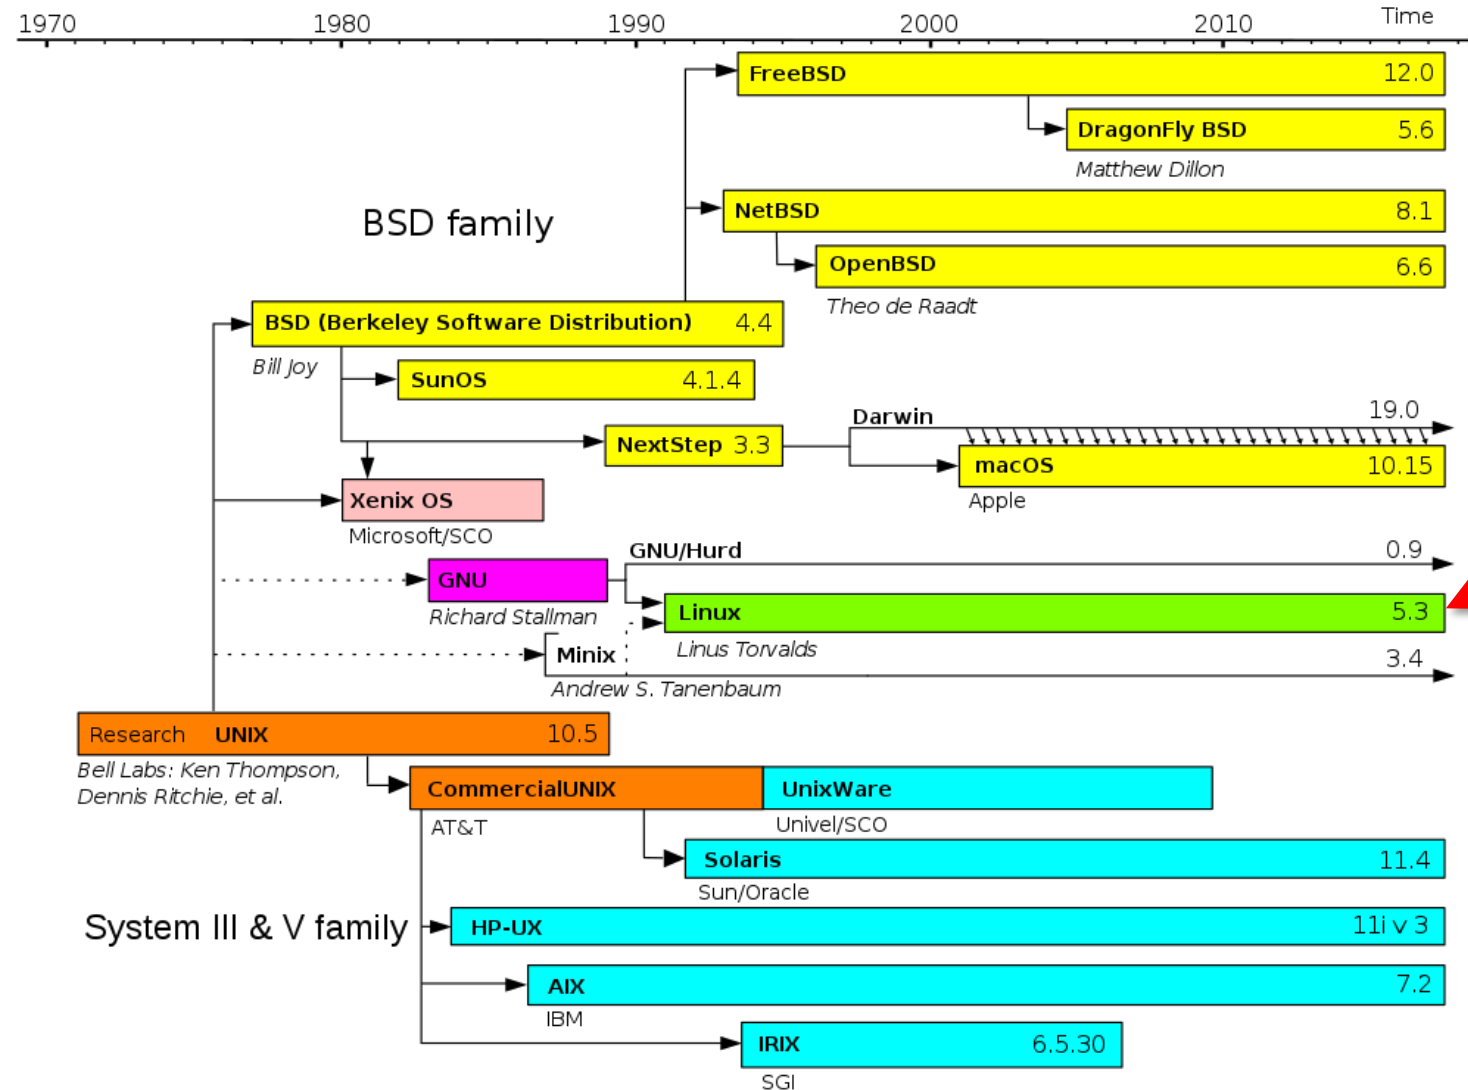
How “Linux” is pronounced?

- /'linʊks/
- Listen at <https://upload.wikimedia.org/wikipedia/commons/0/03/Linus-linux.ogg>
- or if available click the icon below:



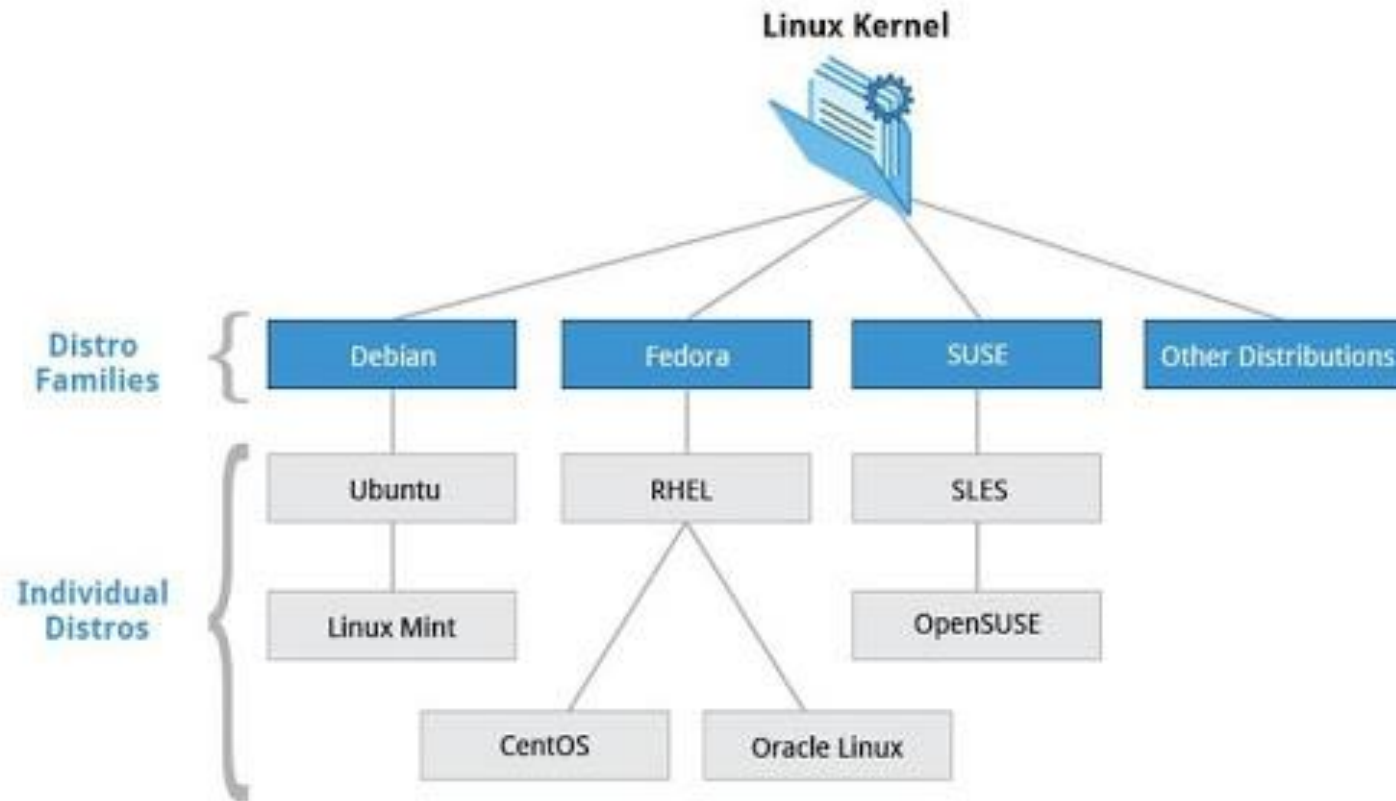
The name "Linux" is also used for a laundry detergent made by Swiss company Rösch.^[190]

Before Linux Timeline



You are here

Some popular linux distribution families



Linux Distribution Timeline

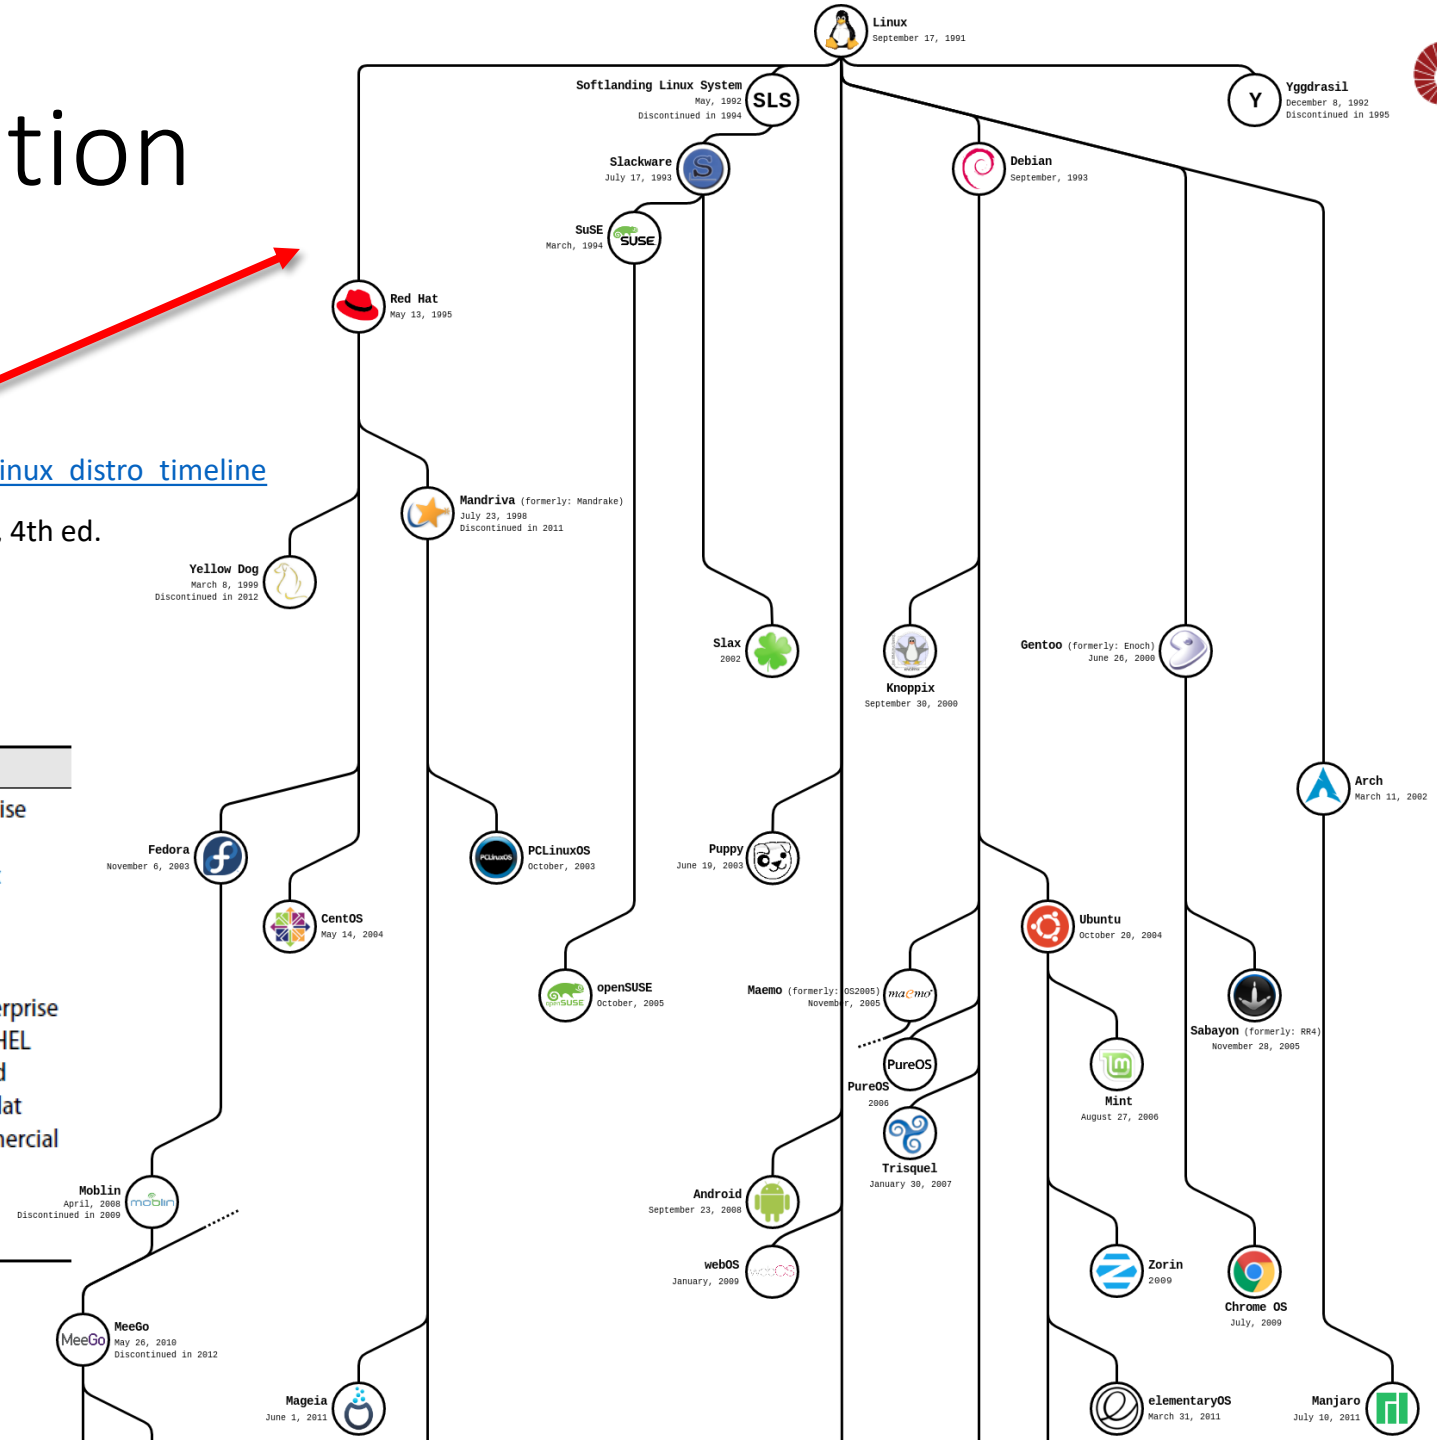
source:

https://commons.wikimedia.org/wiki/Category:Linux_distro_timeline

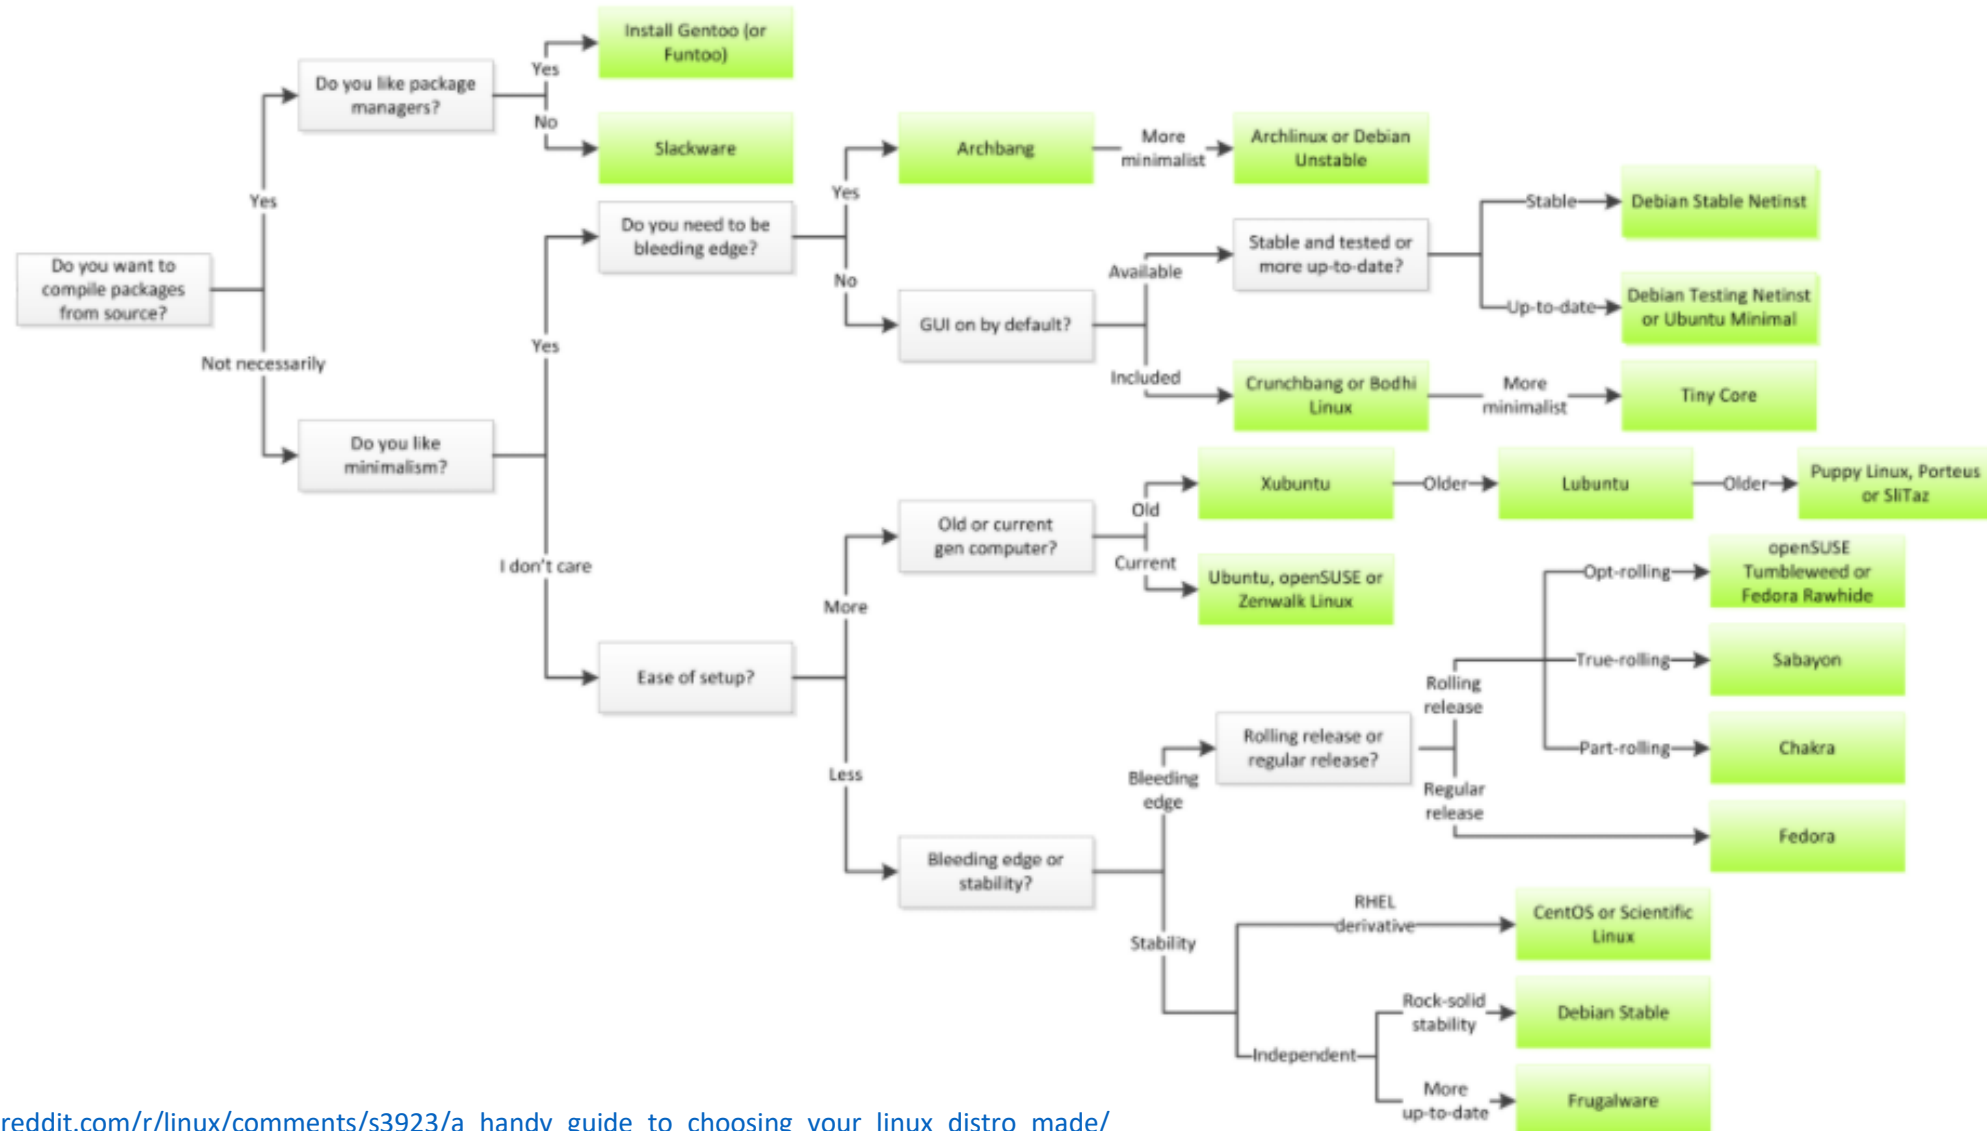
Unix and Linux System Administration Handbook, 4th ed.

Most popular general-purpose Linux distributions

Distribution	Web site	Comments
CentOS	centos.org	Free analog of Red Hat Enterprise
Debian	debian.org	Closest to GNU
Fedora	fedoraproject.org	De-corporatized Red Hat Linux
Gentoo	gentoo.org	Compile-it-yourself, optimized
Linux Mint	linuxmint.com	Ubuntu-based, elegant apps
Mandriva	mandriva.com	Long history, "easy to try"
openSUSE	opensuse.org	Free analog of SUSE Linux Enterprise
Oracle Enterprise Linux	oracle.com	Oracle-supported version of RHEL
PCLinuxOS	pclinuxos.com	Fork of Mandriva, KDE-oriented
Red Flag	redflag-linux.com	Chinese distro, similar to Red Hat
Red Hat Enterprise	redhat.com	Reliable, slow-changing, commercial
Slackware	slackware.com	Grizzled, long-surviving distro
SUSE Linux Enterprise	novell.com/linux	Strong in Europe, multilingual
Ubuntu	ubuntu.com	Cleaned-up version of Debian



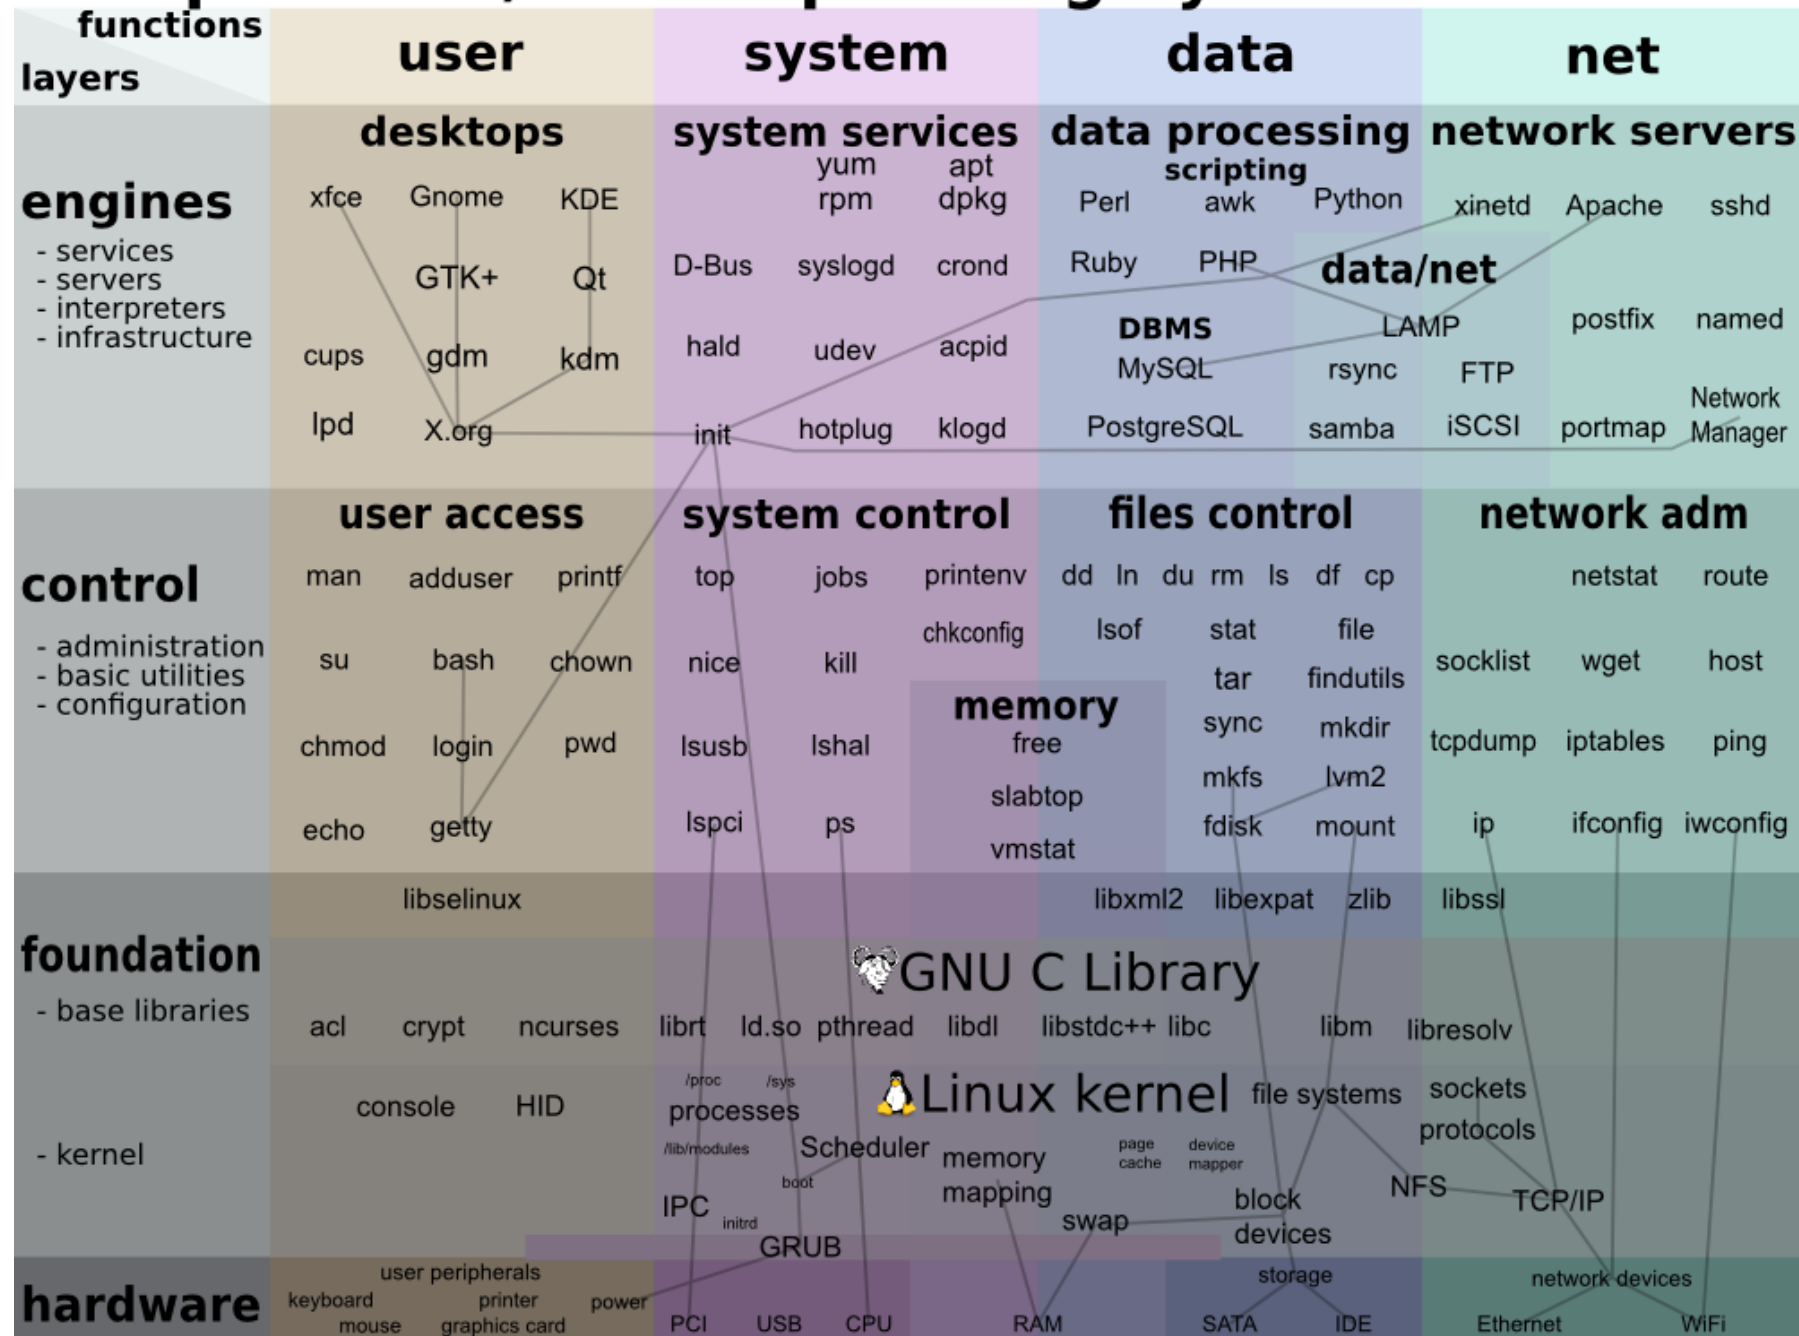
Choose a distro, any distro



Top 10 distros (for personal use) for 2020

POSITION	2020
1	MX Linux
2	Manjaro
3	Linux Mint
4	Ubuntu
5	Debian
6	Elementary OS
7	Solus
8	Zorin OS
9	Fedora
10	Deepin

Map of GNU/Linux Operating System Internals



Practice

- Download and Install Oracle VirtualBox
- Download the latest Ubuntu or Linux Mint Distribution
- Create a virtual machine on VirtualBox
- Install the downloaded OS (do not run as live CD, install to virtual disk)
- Boot the OS
- Open a terminal
- Run a few commands like: ls, cd, mkdir, rm (be careful), man, ps, top, nano



note: you'll be using this virtual machine a lot during the rest of the course.